

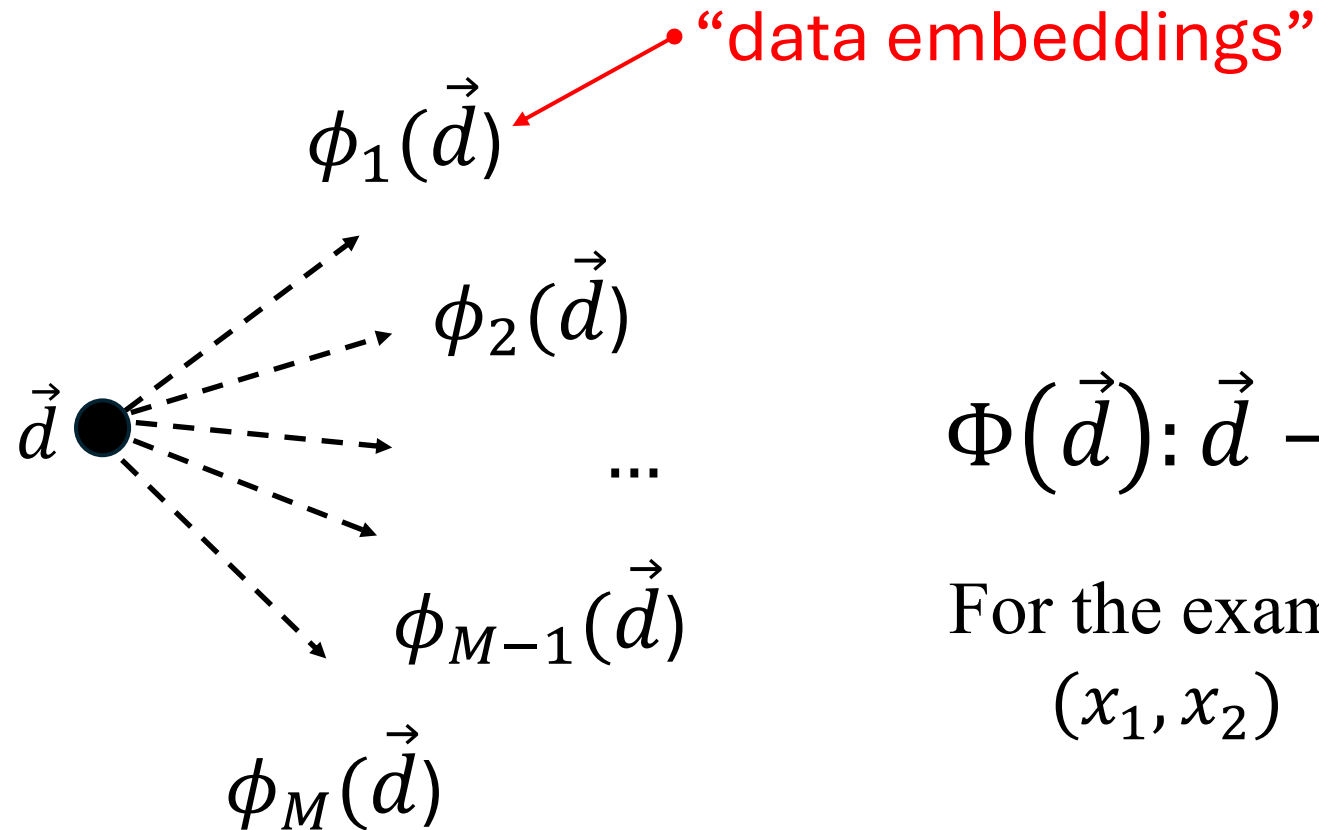
CS 461: Machine Learning Principles

Class 5: Sept. 19

Linear Regression and Data Pre-Processing

Instructor: Diana Kim

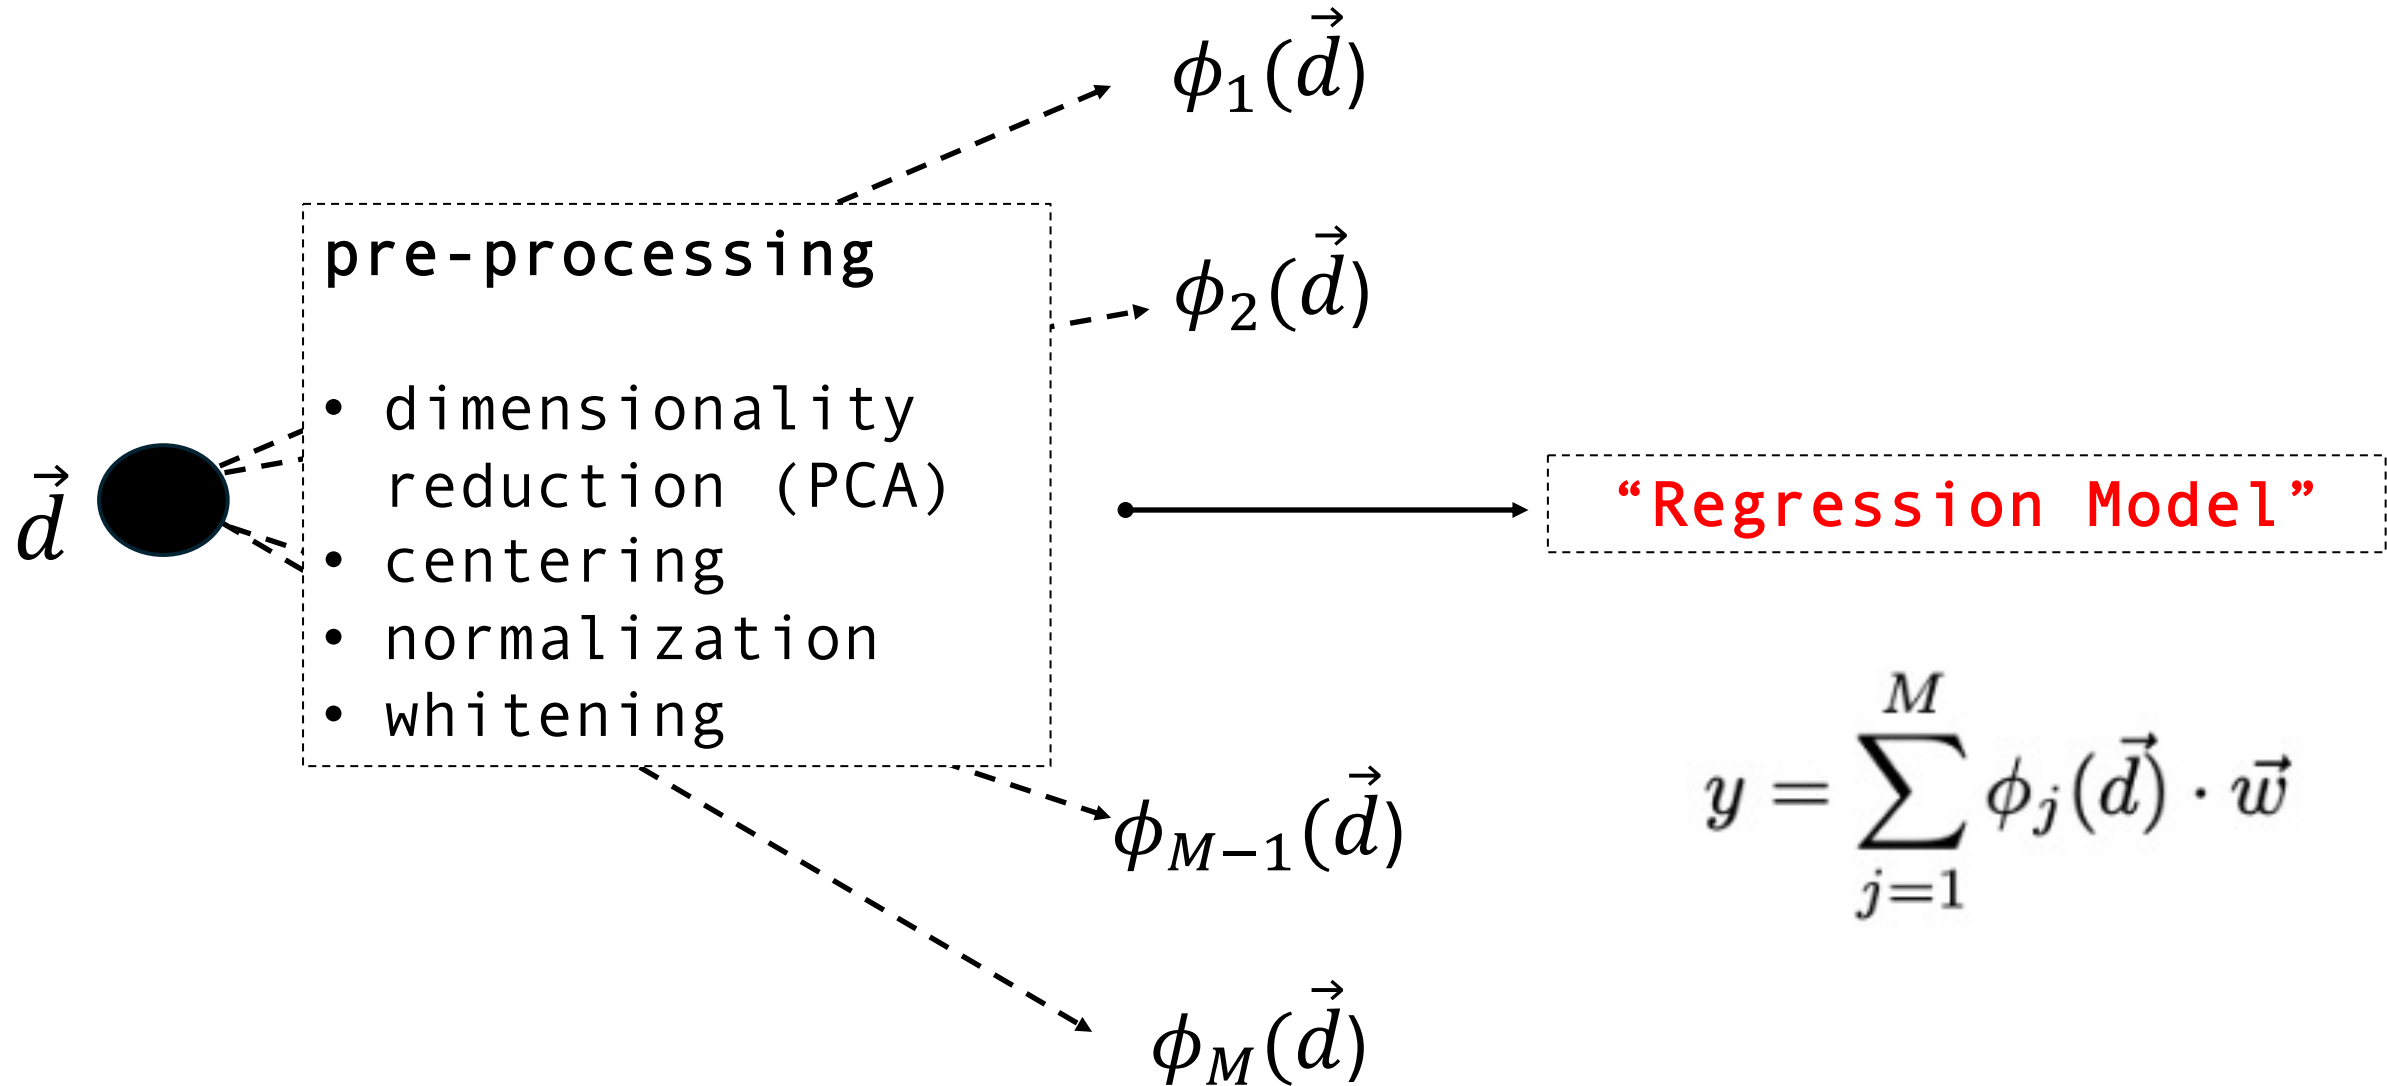
In the last class,
we talked about creating a feature space to accommodate the linear modeling
for regression/ classification problems.



$$\Phi(\vec{d}): \vec{d} \rightarrow R^M$$

For the example of **XOR** problem,
 $(x_1, x_2) \rightarrow (x_1, x_2, x_1 \cdot x_2)$


We may need the pre-processing steps before the algorithm but let's focus on “**regression algorithm**” first and cover pre-processing part.



Regression Problem

Regression Problem

- Suppose we defined a proper feature map $\phi(\vec{d})$ for data point (\vec{d}, y) .
The data matrix D is transformed into Φ .



Regression Problem

- Suppose we defined a proper feature map $\phi(\vec{d})$ for data point (\vec{d}, y) .
The data matrix D is transformed into Φ .

We set the linear combination of $\phi(\vec{d})$ with \vec{w} to predict the value y .

$$y = \Phi(\vec{d}) \cdot \vec{w} + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

ε errors from :

- + imperfection feature space design
- + imperfection hypothesis space
- + error from measurement

Regression Problem

- Suppose we defined a proper feature map $\phi(\vec{d})$ for data point (\vec{d}, y) . then we can transform a data matrix D into Φ .
- We set the linear combination of $\phi(\vec{d})$ with \vec{w} to predict the value y .

$$y = \Phi(\vec{d}) \cdot \vec{w} + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

- We have observed data.
(design matrix Φ)

$\phi_1(d1)$	$\phi_2(d1)$...	$\phi_M(d1)$
$\phi_1(d2)$	$\phi_2(d2)$...	$\phi_M(d2)$
...			
...			
...			
...			
$\phi_1(dN)$	$\phi_2(dN)$...	$\phi_M(dN)$

×

w_1
w_2
...
w_M

=

y_1
y_2
...
y_n

Regression Problem

- Suppose we defined a proper feature map $\phi(\vec{d})$ for data point (\vec{d}, y) . then we can transform a data matrix D into Φ .
- We set the linear combination of $\phi(\vec{d})$ with \vec{w} to predict the value y .

$$y = \Phi(\vec{d}) \cdot \vec{w} + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

- We have observed data.
- We want to estimate \vec{w}

$\phi_1(d1)$	$\phi_2(d1)$...	$\phi_M(d1)$
$\phi_1(d2)$	$\phi_2(d2)$...	$\phi_M(d2)$
...			
...			
...			
...			
$\phi_1(dN)$	$\phi_2(dN)$...	$\phi_M(dN)$

×

w_1
w_2
...
w_M

=

y_1
y_2
...
y_n

Regression Algorithm
: how to find the \vec{w} ?

Regression Problem: Estimation Problem

from observations

$$y = \Phi(\vec{d}) \cdot \vec{w} + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2)$$

want to estimate W !

We learned the two possible ways (MAP and MLE)!

- $w^* = \operatorname{argmax} p(\vec{w}|y, \Phi) = \frac{p(y|\vec{w}, \Phi)p(\vec{w})}{p(\Phi)} : (\text{MAP})$
- $w^* = \operatorname{argmax} p(y|\vec{w}, \Phi) : (\text{MLE})$
- Q: the distribution of y ?

Regression Problem: Estimation Problem

Suppose we have a feature map $\phi(\vec{d})$ for a data point \vec{d}
we want to learn \vec{w} whose the linear combination of $\phi(\vec{d})$ predicts the value y with error $\varepsilon \sim N(0, \sigma^2)$.

we have observations: data Φ ($N \times M$) and \vec{y}

- $w^* = \operatorname{argmax} p(w^* | \vec{y}, \Phi) = \frac{p(y | \vec{w}, \Phi) p(\vec{w})}{p(\Phi)} : (\text{MAP})$
- $w^* = \operatorname{argmax} p(\vec{y} | \vec{w}, \Phi) : (\text{MLE})$
 $= \operatorname{argmax} \mathcal{N}_y(\Phi(\vec{d}) \cdot \vec{w}, \sigma^2 \mathbf{I})$ (when observations are i.i.d)

Regression Problem: Estimation Problem

we have observations: data Φ ($N \times M$) and \vec{y}

- $w^* = \operatorname{argmax} p(\vec{y}|\vec{w}, \Phi)$: (MLE) Ground Truth (data)
 $= \operatorname{argmax} \mathcal{N}_y(\Phi(\vec{d}) \cdot \vec{w}, \sigma^2 \mathbf{I})$

$$\operatorname{argmin}_w ||\vec{y} - \Phi \cdot \vec{w}||^2$$

Prediction!

$$\operatorname{argmin}_w ||\vec{y} - \Phi \cdot \vec{w}||^2$$

MLE becomes
Minimum Mean Square Error Problem

$$J(\vec{w}) = ||\vec{y} - \Phi \cdot \vec{w}||^2$$

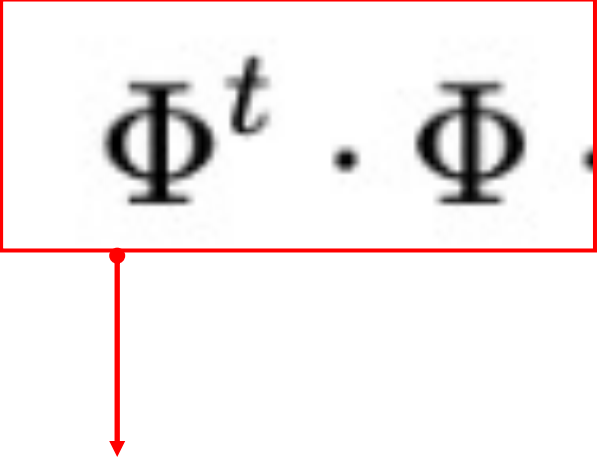
$$J(\vec{w}) = (\vec{y}^t - \vec{w}^t \cdot \Phi^t) \cdot (\vec{y} - \Phi \cdot \vec{w})$$

$$\nabla J(\vec{w}) = -2 \cdot \Phi^t \cdot (\vec{y} - \Phi \cdot \vec{w}) = 0$$

$$\Phi^t \cdot \Phi \cdot \vec{w} = \Phi^t \cdot \vec{y}$$

Normal Equation¹³

Regression Problem: Estimation Problem

$$\Phi^t \cdot \Phi \cdot \vec{w} = \Phi^t \cdot \vec{y}$$


The three possible cases:

- invertible (Rank M)
- invertible (Rank M) but close to singular (very small eigenvalues)
- non – invertible (Rank $< M$)

Note) close look into the bias term

$$J(w) = ||\vec{y} - \Phi \cdot \vec{w}'||^2$$

$$= ||\vec{y} - \Phi' \cdot \vec{w}' - [b, b, \dots b]^t||^2$$

$$\frac{\partial J}{\partial b} = -2 \cdot [1, 1, 1, \dots, 1] \cdot (\vec{y} - \Phi' \cdot \vec{w}' - [b, b, \dots b]^t) = 0$$

$$= \sum_{i=1}^N y_i - \sum_{i=1}^N y'_i - b \cdot N = 0$$

$$bN = \sum_{i=1}^N y_i - \sum_{i=1}^N y'_i$$

$$b = \frac{1}{N} \sum_{i=1}^N y_i - \frac{1}{N} \sum_{i=1}^N y'_i$$

ϕ_1 (d1)	ϕ_2 (d1)	...	1
ϕ_1 (d2)	ϕ_2 (d2)	...	1
...			
...			
...			
...			
ϕ_1 (dN)	ϕ_2 (dN)	...	1

×


W_1
W_2
...
b

=

y_1
y_2
...
y_n

- $y' = \text{prediction without bias?}$

Geometrical Interpretation: MMSE Solution



- + The overdetermined system $\Phi w = y$ might not have a solution, so we find an approximated solution instead.
- + $\Phi^t \Phi w = \Phi^t y$ The approximated solution is the MMSE solution. MMSE can be derived from MLE.

Geometrical Interpretation MMSE Estimation

$$\Phi(\vec{d}) \cdot \vec{w} = \vec{y}$$

- solving overdetermined System

- projection to the data space and find an approximated solution

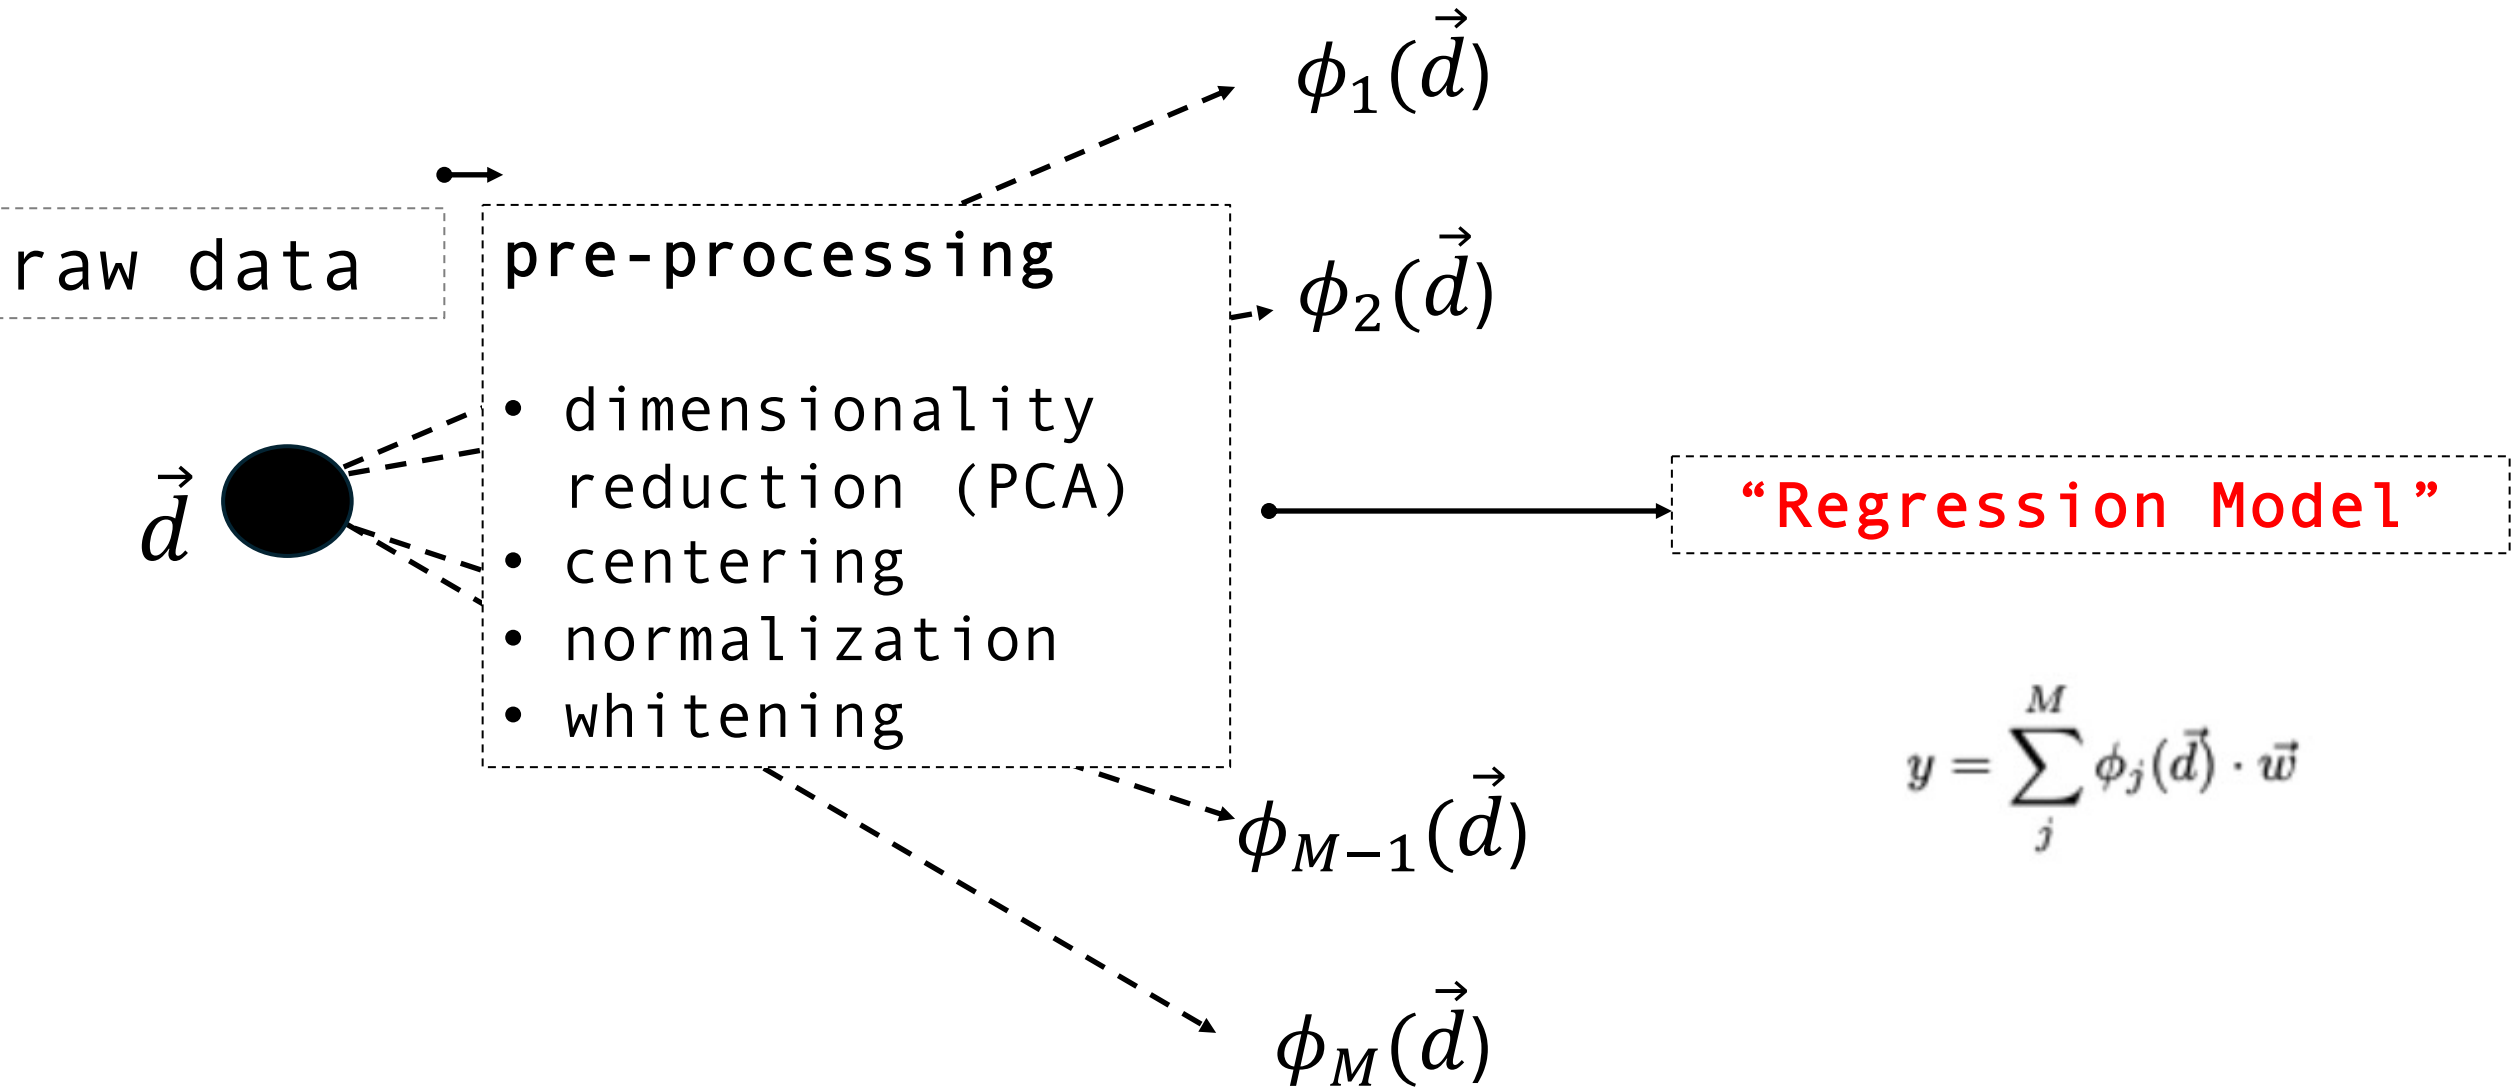
$$\Phi(\vec{d})^t \cdot \Phi(\vec{d}) \cdot \vec{w} = \Phi(\vec{d})^t \cdot \vec{y}$$

- when data space's rank is M : unique approximated solution

- when data space's rank is less than M: pseudo-inverse solution

We learned regression algorithm.
but we need some pre-processing steps for successful learning.

We need the pre-processing steps before the algorithm.



$$y = \sum_j^M \phi_j(\vec{d}) \cdot \vec{w}$$

+ for regression,
we will focus on dimensionality reduction and whitening

Two Kinds of Raw Data in Regression Problems

- each feature dimension has semantic.

“prediction of house market price”

single house/ townhome, square feet, garden size, public school scores

- no semantic, a whole vector represents an image/ audio /texts

Both cases will need pre-processing steps, but


Semantic Data

[A car's fuel efficiency in miles per gallon]

Pounds in 1000s (features)	Miles Per Gallon (Label)
3.5	18
3.69	15
3.44	18
3.43	16
4.34	15
4.42	14
2.37	24

From <https://developers.google.com/machine-learning/crash-course/linear-regression>

Regression Model: Prediction of Miles/Gallon



From <https://developers.google.com/machine-learning/crash-course/linear-regression>

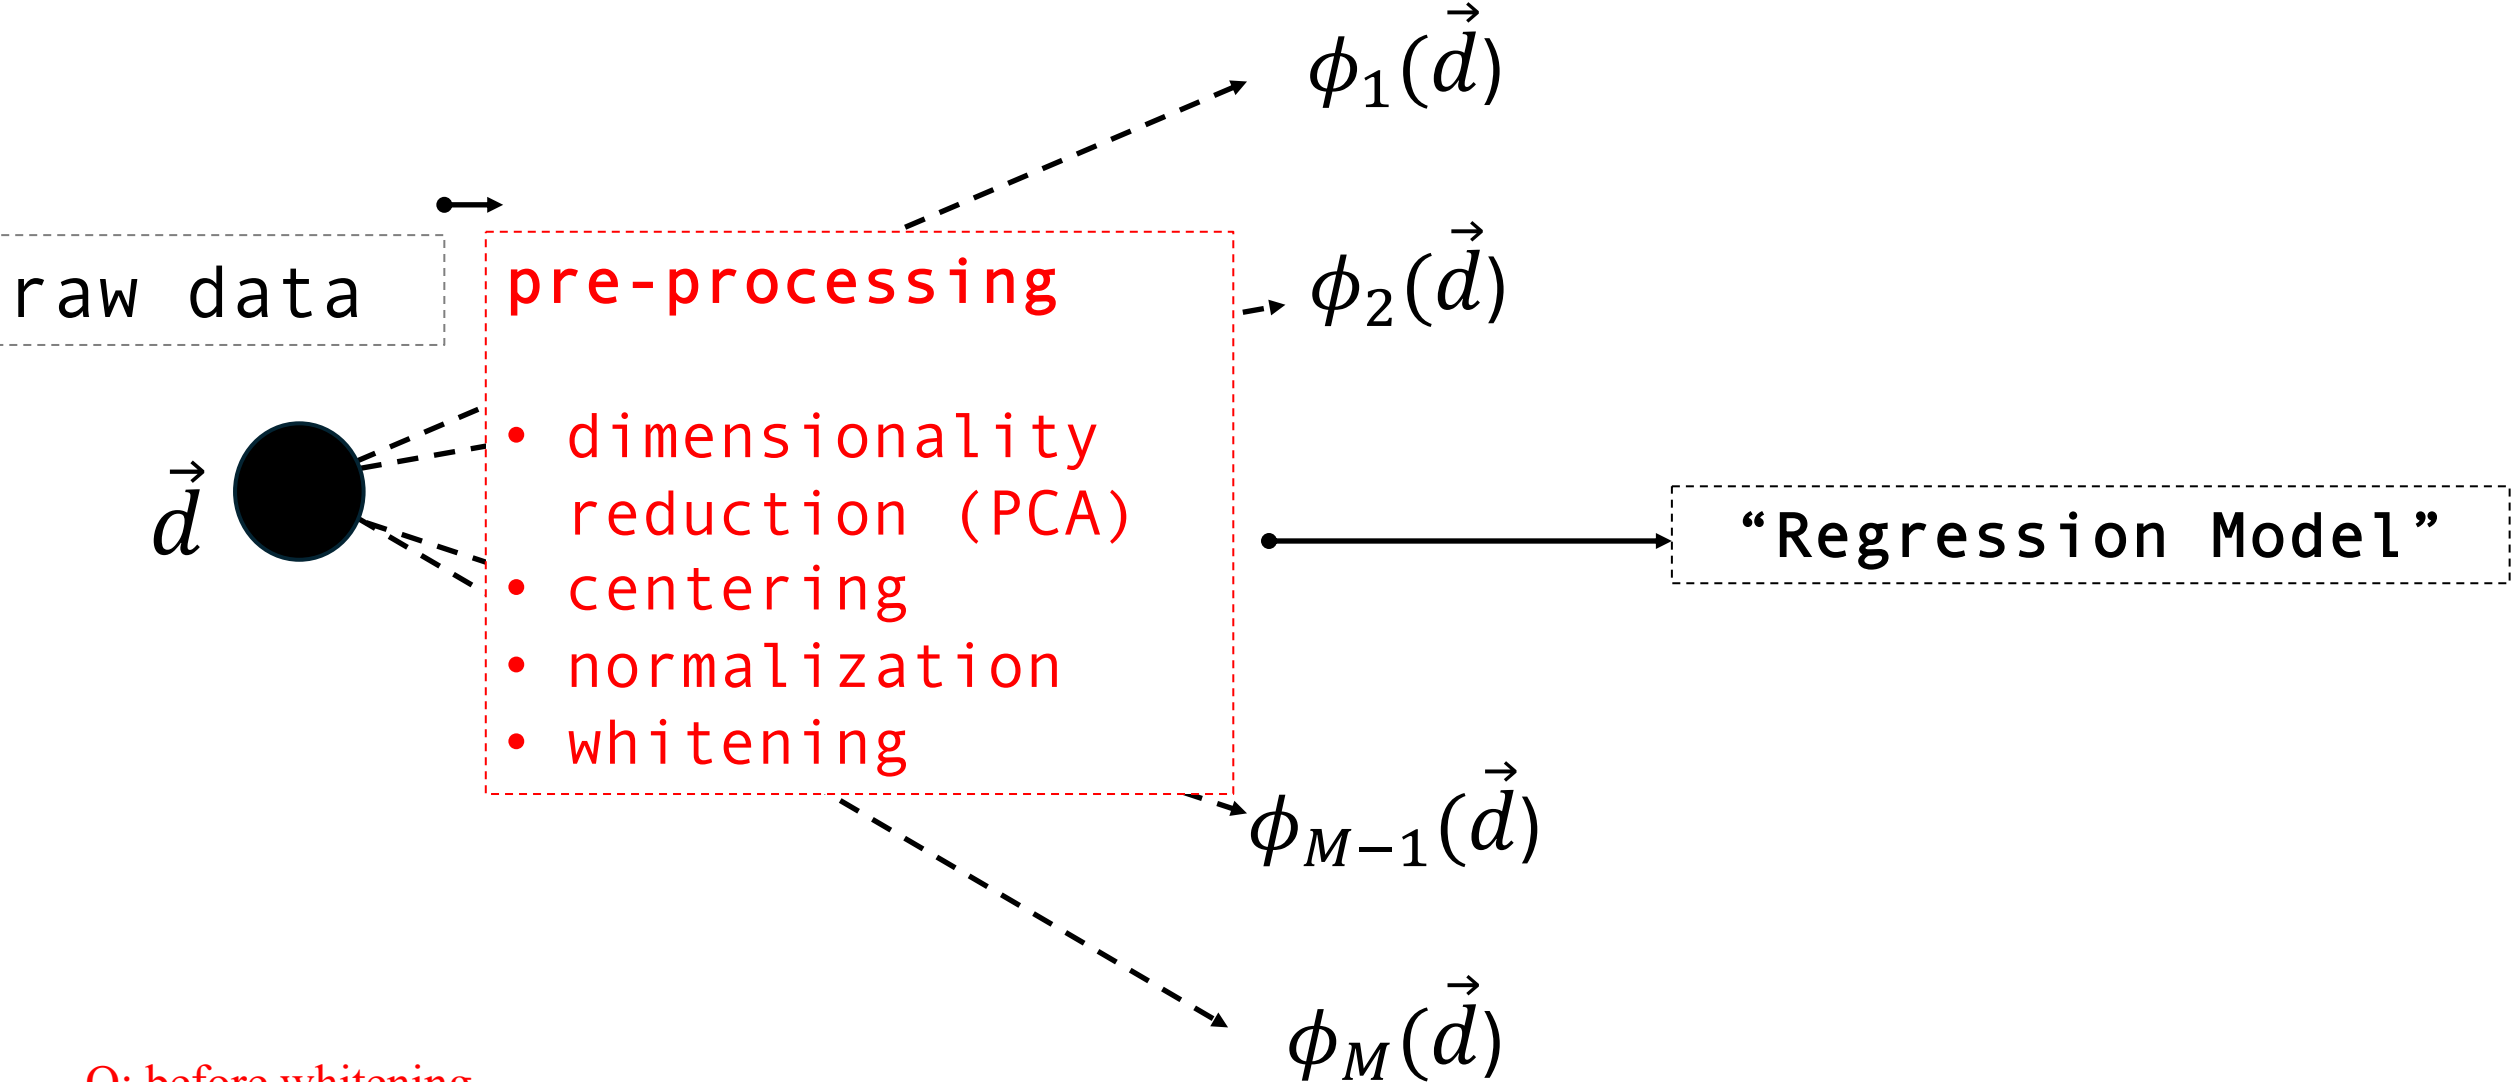
Vector Data (in general very high dimensional)

Image representation in the last hidden layer of deep-CNN

Resurrection, Alma Thomas

1.66116878e-02	1.43487025e-02	2.32550185e-02	1.19144898e-02
9.63837430e-02	8.31691474e-02	3.06324158e-02	2.04893108e-02
3.52624245e-02	5.56392968e-03	1.73700173e-04	1.50486574e-01
1.23705138e-02	5.22252880e-02	5.49408374e-03	4.67238687e-02
1.44064635e-01	5.96226342e-02	3.30738463e-02	7.86291659e-02
1.78842451e-02	8.52992311e-02	1.89456772e-02	1.03305764e-02
3.74222398e-02	8.35433230e-02	1.95080508e-02	3.26446295e-02
2.95458261e-02	1.08750183e-02	1.02128655e-01	4.72795311e-03
8.63977734e-02	2.39002742e-02	1.02527821e-02	2.37334445e-02
1.44393004e-02	2.45263092e-02	4.37902249e-02	2.01957620e-04
1.88012738e-02	3.28723639e-02	2.84272023e-02	9.55437776e-03
4.14613076e-02	1.39710018e-02	1.36978943e-02	1.87548213e-02
8.24719146e-02	4.43845317e-02	3.53335054e-03	7.63716456e-03
9.22968891e-03	3.64266671e-02	2.88719591e-02	1.30119538e-02
2.02936288e-02	1.33205568e-02	1.15464339e-02	8.20994973e-02
1.77106280e-02	8.52256175e-03	1.17111830e-02	1.45943370e-02
9.04859081e-02	4.89737056e-02	4.16163765e-02	1.04203597e-02
2.58007385e-02	4.69408296e-02	4.55647372e-02	1.24655450e-02
2.41902079e-02	2.42145211e-02	4.81210562e-04	8.90940428e-03
6.85443357e-02	4.78595048e-02	1.18027581e-02	1.17037995e-02
1.90981105e-02	1.00829145e-02	5.23220561e-03	3.84746492e-02
8.81355628e-02	3.36198583e-02	5.35092168e-02	4.87123579e-02
8.99140537e-03	5.39787523e-02	4.79393527e-02	2.99579669e-02
1.71675645e-02	3.76332477e-02	7.88647458e-02	3.79528590e-02
4.51750029e-03	9.38767791e-02	3.61216962e-02	1.98117495e-02
3.94294709e-02	1.14793226e-01	1.48017062e-02	6.40132884e-03
8.16167146e-03	6.94637448e-02	3.43800858e-02	8.04584008e-03
1.55022442e-01	2.59600277e-03	3.20520252e-02	1.00370266e-01
6.82575488e-03	3.21909995e-03	6.07831627e-02	5.22131985e-03
4.10482734e-02	5.29111736e-02	4.37461957e-02	4.37461808e-02
4.10865359e-02	9.59158130e-03	4.68185917e-02	7.04082549e-02
5.19240461e-03	9.47480425e-02	1.72703192e-02	1.32609099e-01
1.84857957e-02	2.34019849e-03	2.21508313e-02	3.19128227e-03
1.03731174e-02	7.90489465e-02	3.40001471e-02	2.08658073e-02
3.63909267e-03	2.93061193e-02	1.79619715e-02	3.92507110e-03
1.22312911e-01	4.27385271e-02	4.02529091e-02	6.87315594e-03
1.79619640e-02	1.44496362e-03	3.47868539e-04	2.03075245e-01
2.45202169e-01	1.26138151e-01	1.07999377e-01	1.46901429e-01
9.70007405e-02	1.03836969e-01	1.09804377e-01	1.04106106e-01
8.70869756e-02	8.81577432e-02	7.79228508e-02	9.59928930e-02
2.06121951e-01	2.38734394e-01	1.37491360e-01	7.11895898e-02
9.10348147e-02	1.08147562e-01	8.93435627e-02	8.45326930e-02
8.54639262e-02	7.94288218e-02	7.84831643e-02	6.98279142e-02
6.67123348e-02	7.02826679e-02	1.02719694e-01	1.04542613e-01
1.12103589e-01	8.02482218e-02	1.26211137e-01	1.22251317e-01
1.18328705e-01	9.65996012e-02	9.47735459e-02	8.21543038e-02
7.41177499e-02	1.03439212e-01	1.14290312e-01	1.15447372e-01
1.28355548e-01	1.06327742e-01	7.30694234e-02	6.20305464e-02
1.06132567e-01	7.94187784e-02	8.86070132e-02	8.47868249e-02
1.07920051e-01	8.36525112e-02	6.55624866e-02	7.80229717e-02
8.64467472e-02	8.55527893e-02	1.10759147e-01	1.32106051e-01
7.44441077e-02	5.27140088e-02	1.08958408e-01	8.06024447e-02
9.61078107e-02	9.56790447e-02	1.04670644e-01	8.01085979e-02
6.94930553e-02	7.93032944e-02	9.49410051e-02	7.71025643e-02
1.05781302e-01	1.46627113e-01	6.05126023e-02	4.13953587e-02
1.23981662e-01	1.08231083e-01	1.34232193e-01	1.18365087e-01
1.09341882e-01	8.85261148e-02	8.09772611e-02	8.30637813e-02
1.13967851e-01	8.66363198e-02	1.12511240e-01	1.40123367e-01
6.65690452e-02	4.43194956e-02	1.57082587e-01	1.04566351e-01
9.03969407e-02	1.14839226e-01	1.21048279e-01	9.68690664e-02
9.22555625e-02	1.10619672e-01	1.21558294e-01	8.79304186e-02
1.04587585e-01	1.42198473e-01	8.80973265e-02	4.42507043e-02

We need the pre-processing steps for the raw data.

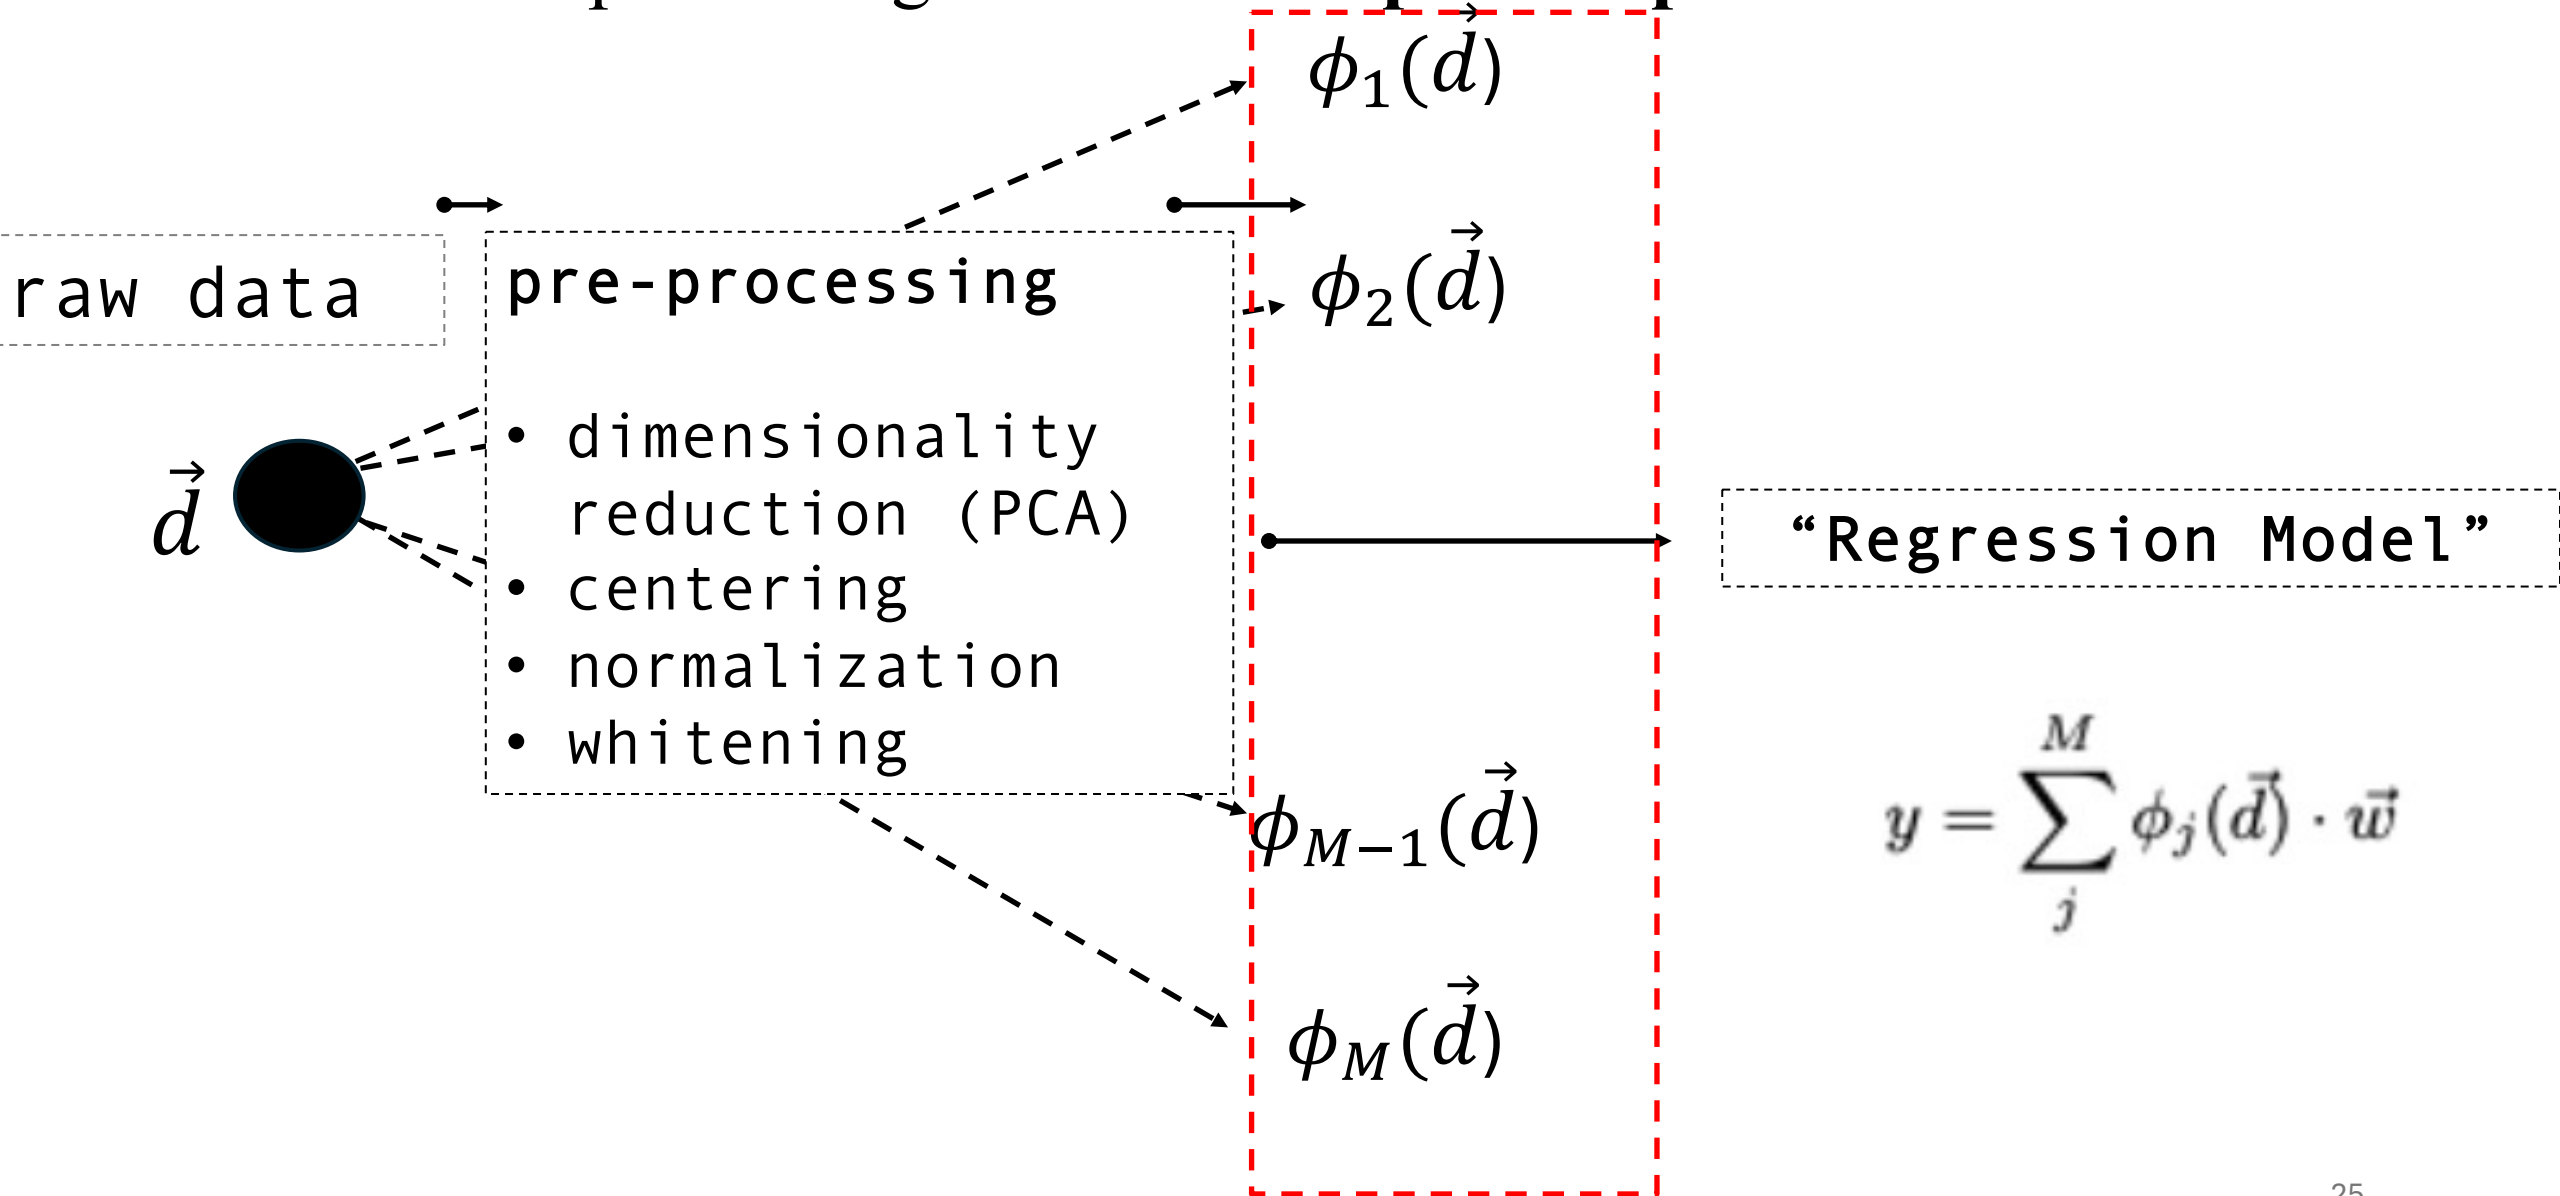


Q: before whitening,
We must do PCA. Why? Otherwise, some error dimension can be magnified by whitening (1/small eigenvalue)

Centering, Normalization, Standardization, Whitening


- Centering: $\vec{x} - E[\vec{X}]$
- Normalization: $\frac{\vec{x}}{||x||}$
- Standardization: $\frac{x_i - E[X_i]}{VAR[X_i]}$ (element-wise operation)
- Whitening : $Y = AX + \vec{b}$ when $Y \sim N(0, I)$

Raw data \longrightarrow Pre-processing \longrightarrow **Feature Space Expansion with Basis**



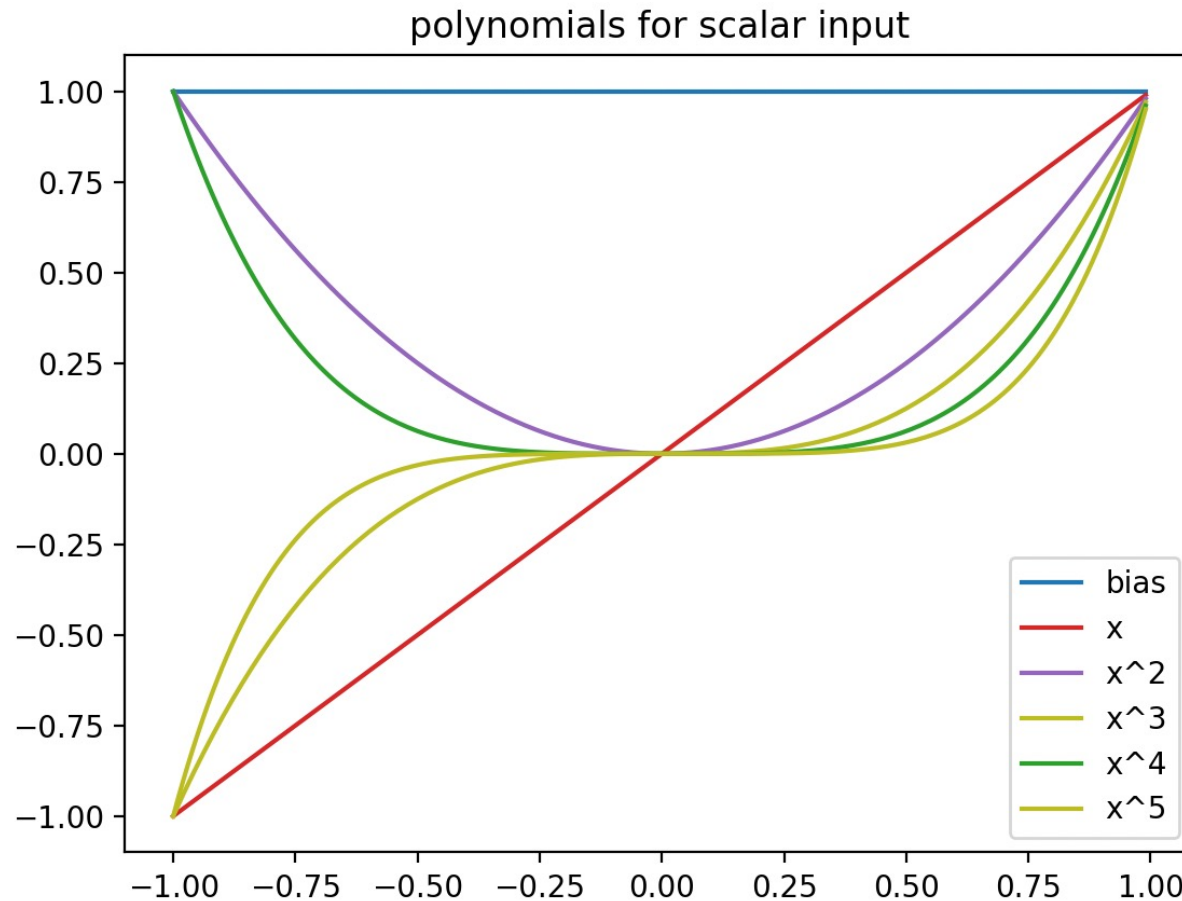
Basis Function (1) : Polynomial Expansion

+ Pre-processed data:
(reduced dimensions and whitened)



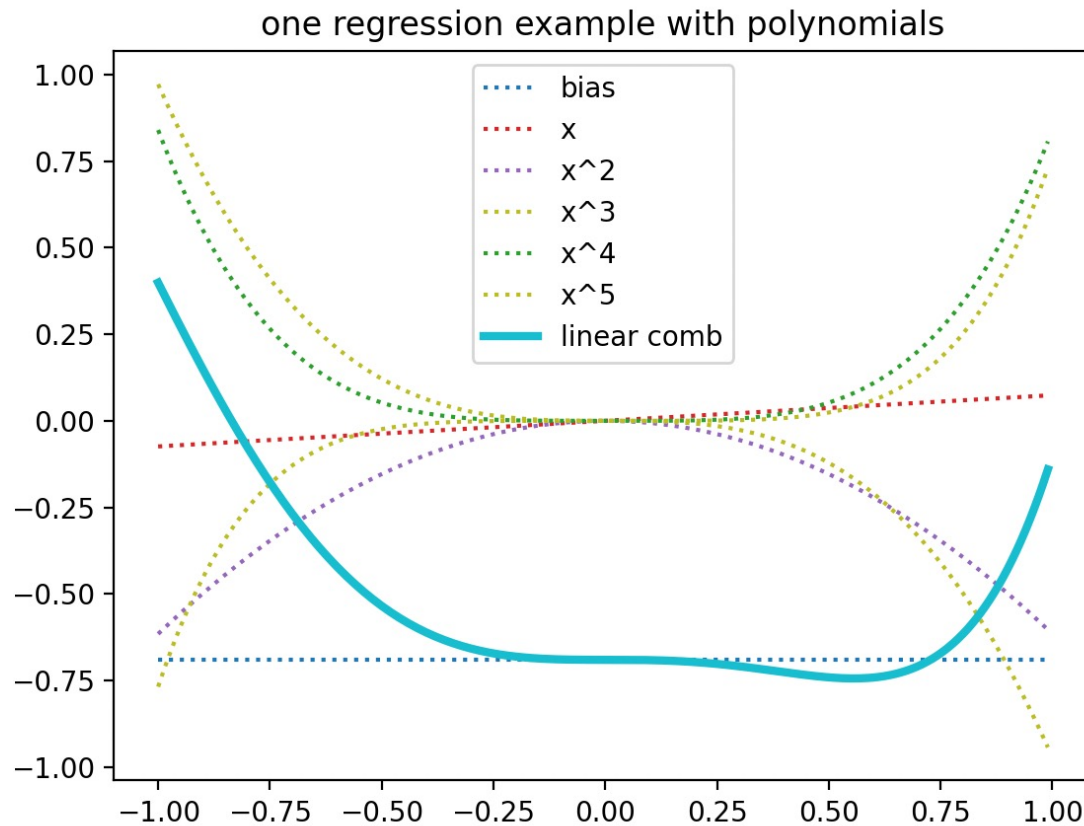
	1	x_2	x_2^2
1	1	x_2	x_2^2
x_1	x_1	x_1x_2	\times
x_1^2	x_1^2	\times	\times

Basis Function (1) : Polynomial Expansion



- This example shows the case when input is scalar.
- Q: what if input is a 2D vector? How would you draw the plot?

Basis Function (1) : Polynomial Expansion



- This example shows the case when we set \vec{w} with an arbitrary parameter. This shows a possible regression function.

Basis Function (2) : Gaussian Basis function



$$\phi_j = \exp \left\{ -\frac{(x - \mu_j)^2}{2\sigma^2} \right\}$$

Q: the magnitude of σ^2 ?

(small: local and spiky vs. large: global and smooth)

Q: the locations of μ_j ? (dense / sparse)

Basis Function (2) : Gaussian Basis function



The magnitude of σ determines the influence over other neighboring Gaussian functions.

Summary: **Train** and **Test** Procedures for Regression Problem

- A set of hypothetical basis functions (polynomial / Gaussian)
- We have data points $((d_1, y_1), (d_2, y_2), (d_3, y_3), \dots, (d_N, y_N))$

Train

- Train data PCA for dim-reduction (high-dimensional data) and whitening. (save the PCA blocks computed with training set)
- Normal equation to estimate W

Test

- Test data PCA for dim-reduction and whitening with the same block used in training!!!! (important)
- Compute error between the ground truth y and prediction y' : $\|y - y'\|^2$


Dimensionality Reduction: PCA

- how can we reduce high-dimensional images into 2D space?

Principal Component Analysis (PCA)

We want to make data (N-D) moves around within a subspace (M-D), ($N > M$)

Vector $\vec{x} = (a, b)$ can be represented by different orthonormal bases.



When $x_n \in R^D$ and u_i where $i = 1, 2, \dots, D$ are are abnormal basis,

- X_n

$$x_n = \sum_{i=1}^D \alpha_{ni} u_i \quad \text{and} \quad \alpha_{ni} = \langle x_n, u_i \rangle$$

- \widetilde{X}_n

We want to approximate x_n on the subspace of the first M basis,

$$\tilde{x}_n = \sum_{i=1}^M z_{ni} u_i + \sum_{i=M+1}^D b_i u_i$$

Q: What is the optimal values for Z_{ni} and b_i minimizing $J = \| X_n - \widetilde{X}_n \|^2$

We want to minimize the averaged square error between x_n and \tilde{x}_n

$$\arg \min_{(z_{ni}, b_i)} J = \frac{1}{N} \sum_{n=1}^N \left(\langle x_n, u_i \rangle \cdot u_i^t - \sum_{i=1}^M z_{ni} u_i^t - \sum_{i=M+1}^D b_i u_i^t \right) \cdot$$

$$\left(\langle x_n, u_i \rangle \cdot u_i - \sum_{i=1}^M z_{ni} u_i - \sum_{i=M+1}^D b_i u_i \right)$$

- Respect to z_{nk}

$$\frac{\partial J}{\partial z_{nk}} = (-2u_k^t) \cdot \left(\langle x_n, u_i \rangle \cdot u_i - \sum_{i=1}^M z_{ni} u_i - \sum_{i=M+1}^D b_i u_i \right)$$

$$= -2 \langle x_n, u_k \rangle + 2z_{nk} = 0$$

- Respect to b_r

$$\frac{\partial J}{\partial b_r} = \frac{1}{N} \sum_{n=1}^N (-2u_r^t) \cdot \left(\langle x_n, u_i \rangle \cdot u_i - \sum_{i=1}^M z_{ni} u_i - \sum_{i=M+1}^D b_i u_i \right)$$

$$= \frac{1}{N} \sum_{n=1}^N (-2 \langle x_n, u_r \rangle + 2b_r)$$

- Rewrite J

$$J = ||x_n - \tilde{x}_n||^2 = \frac{1}{N} \sum_{n=1}^N \sum_{i=M+1}^D ((x_n - \bar{x})^t u_i)^2$$

$$= \sum_{i=M+1}^D u_i^t \cdot \frac{1}{N} \sum_{n=1}^N (x_n - \bar{x})(x_n - \bar{x})^t \cdot u_i$$

$$= \sum_{i=M+1}^D u_i^t \Sigma u_i$$

Estimation of COV (X,X)

- Lagrangian Function for the Constraint $u_i^t u_i = 1$

Lagrangian function for the constraint $\|u_i\| = 1$

$$J(\lambda) = \sum_{i=M+1}^D u_i^t \Sigma u_i + \lambda(1 - u_i^t u_i)$$
$$\frac{\partial J}{\partial u_i} = \Sigma u_i - \lambda^* u_i = 0$$

Q: What the optimal solution indicate about u_i ?

- Go back to J

$$J = ||x_n - \tilde{x}_n||^2 = \frac{1}{N} \sum_{n=1}^N \sum_{i=M+1}^D ((x_n - \bar{x})^t u_i)^2$$

$$= \sum_{i=M+1}^D u_i^t \cdot \frac{1}{N} \sum_{n=1}^N (x_n - \bar{x})(x_n - \bar{x})^t \cdot u_i$$

$$= \sum_{i=M+1}^D u_i^t \Sigma u_i$$

Estimation of COV (X,X)

Q: To minimize J?

Now, we are ready to define \widetilde{X}_n (*PCA Approximation*)

$$\begin{aligned}\tilde{x}_n &= \sum_{i=1}^M (x_n^t u_i) u_i + \sum_{i=M+1}^D (\bar{x}^t u_i) u_i \\ &= \bar{x} - \bar{x} + \sum_{i=1}^M (x_n^t u_i) u_i + \sum_{i=M+1}^D (\bar{x}^t u_i) u_i \\ &= \bar{x} - \sum_{i=1}^M (\bar{x}^t u_i) u_i + \sum_{i=1}^M (x_n^t u_i) u_i \\ &= \bar{x} + \sum_{i=1}^M ((x_n^t - \bar{x}^t) u_i) u_i \\ &= \bar{x} + U_M U_M^t (x_n - \bar{x})\end{aligned}$$

\widetilde{X}_n is not full dimension.

Depending on how we select U_M , we can define different approximations.

- Variance of \widetilde{x}_n

$$\widetilde{x}_n - \bar{x} = u_j^t (x_n - \bar{x}) u_j$$

$$\frac{1}{N} (\widetilde{x}_n - \bar{x})^t (\widetilde{x}_n - \bar{x}) = \frac{1}{N} u_j^t (x_n - \bar{x}) u_j u_j^t (x_n - \bar{x})$$

$$\text{var}(\widetilde{x}_n) = \lambda_i$$

Different PCA Approximation for $M = 1$, $M = 10$, $M = 50$, $M = 250$




Figure 12.5 An original example from the off-line digits data set together with its PCA reconstructions obtained by retaining M principal components for various values of M . As M increases the reconstruction becomes more accurate and would become perfect when $M = D = 28 \times 28 = 784$.

From Bishop Chap. 12

$$\widetilde{X}_n = \bar{x} + U_M U_M^t (x_n - \bar{x})$$

Visualization of Mean and Eigenvectors

The image can be represented by sum of mean and the linear combinations of eigenvectors




Figure 12.3 The mean vector \bar{x} along with the first four PCA eigenvectors u_1, \dots, u_4 for the off-line digits data set, together with the corresponding eigenvalues.

From Bishop Chap. 12

$$\widetilde{X}_n = \bar{x} + \underbrace{U_M U_M^t (x_n - \bar{x})}_{\text{A vector}}$$

The linear combination of eigenvectors

PCA Applications

- Compression (small variance dimension does not help in learning)


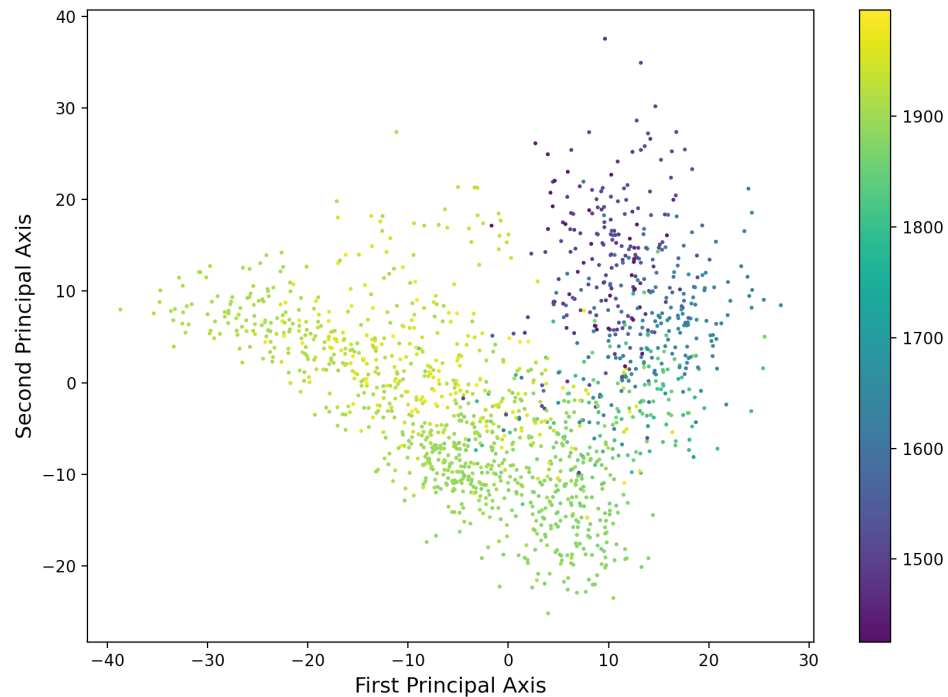
$$\widetilde{X}_n = \bar{x} + U_M U_M^t (x_n - \bar{x})$$

- Whitening (Rotation)

$$\tilde{x}_n = \Lambda^{-\frac{1}{2}} U_M^t (x_n - \bar{x})$$

PCA Applications

- Visualization (1) (the projection of high dimensional data to 3D or 2D)



The last hidden layer embedding of a Deep-CNN Style Classifier is projected to the top principal axes (the eigenvectors corresponding to the first and second largest eigenvalues). The samples are color-coded by year of made.

PCA Applications

- Visualization (2) (projection of high dimensional data to 3D or 2D)




Figure 14. 6 from Deep Learning by Ian Goodfellow

One dimensional manifold that traces out a curved path for vertical shift of digit “9”. The manifold in the high dimensional space is projected into 2D.